

# An Experimental Study of Online Database Index Selection Using MTS

COMP90055 - Research Project Presentation

Tanzid Sultan

# The Need for Online Index Selection

- The simple reason why indexes are important in a database:

```
SELECT * FROM Employees  
WHERE City = 'x';
```

Employees Table

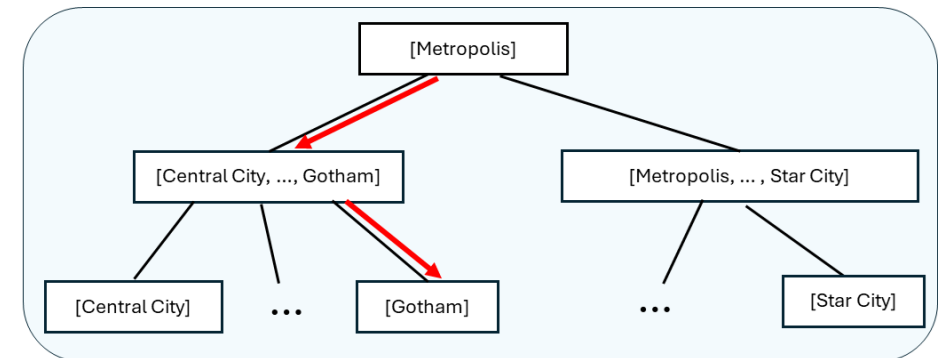
ID	Name	Age	City	Occupation
1	Alice	30	Gotham	Data Analyst
2	Bob	22	Metropolis	Grave Digger
3	Charlie	51	Gotham	Product Manager
...	...	...	...	...

# of rows in table =  $10^9$   $\Rightarrow$  **Full Table Scan = 10 minutes**

# of rows matching predicate = 120  $\Rightarrow$  **Index Scan = 0.01 seconds**

**1000x speedup** with index scan!

B+ Tree Index on City Column



# The Need for Online Index Selection

- Just create an index on every column of every table?
  - Don't have infinite memory and time! (specially when database already large and complex)
- Also, what about **dynamic** workloads and system environments?
  - a static index configuration won't be best all the time (need to adapt)
  - memory budget might vary over time (need to adapt)
  - hand selection by human admin is impractical (need to automate)
- **Solution** → An **Online Algorithm** which **automatically** creates and drops indexes based on evolving workloads and system needs




# Online Index Selection Problem Formulation and MTS

- Queries  $q_i$  arrive sequentially  $\rightarrow$  at each step algorithm decides on changing the index configuration (creating/dropping indexes)  $\rightarrow$  then executes  $q_i$

- Objective is to minimize the **total workload cost**:

$$C_{tot}(\text{Workload}) = \sum_{i=1}^n C_{tr}(s_{i-1}, s_i) + C_{exe}(q_i, s_i)$$

  
*transition cost*                      *execution cost*

- Decision has to be made using **limited knowledge**: only workload/queries seen so far
- One-to-One correspondence with the problem of **metrical task systems** ([Borodin et. al. 1992](#)):

Task  $\leftrightarrow$  Query

State  $\leftrightarrow$  Index Configuration

# WFA and WFIT for Online Index Selection

- The **Work Function Algorithm (WFA)** is an **optimal deterministic algorithm** for MTS ([Borodin et. al. 1992](#)), i.e. achieves lowest possible competitive ratio upper-bound.
- WFA *mimics the* optimal offline algorithm but uses only limited information.

- Should we use WFA for online index selection?

Pros	Cons
<ul style="list-style-type: none"><li>• Simple algorithm</li><li>• Strong performance guarantee</li></ul>	<ul style="list-style-type: none"><li>• Exponentially large configuration space</li><li>• Computationally intractable</li></ul>

- Stable partitions to the Rescue! **WFIT- Divide and Conquer** ([Schnaitter and Polyzotis, 2012](#))

# A Major Flaw of WFIT & Our Research Goals

- A **major weak-point** in WFIT → uses internal cost model

(Internal cost model means cost model of the DB query optimizer, accessed through ***what-if interface***.)

- What's wrong with using *what-if*?
  - query optimizer's **cost estimates are highly error-prone**
  - high overhead of **invoking what-if interface makes WFIT too slow**
- Research goals:
  0. Leverage WFIT for online index selection
  1. **Extend WFIT - replace what-if with a lightweight external cost model**
  2. **Benchmark WFIT against current state-of-the-art MAB**

# A light-weight external cost model for access path prediction

- Difficulty in training **end-to-end ML cost model** (feature engineering, sample inefficiency, model complexity, offline training, etc.)
- Easier to focus directly on **access path prediction** → more relevant for index selection in large disk-based databases.

- Start with **simple heuristic-based approach**:

Input: *query, configuration*

**for each table:**

- **enumerate all possible access paths** (i.e. full table scan + possible index scans)
- **estimate the disk I/O cost for each access path** (**heuristic-based**)
- **pick cheapest access path**

- This **works surprisingly well!** (comparable to *what-if* in accuracy + 500x faster)

# Heuristic-based Selectivity Estimation

- Disk I/O cost for index scans derived using **selectivity estimates**, i.e. **fraction of rows** that need to be retrieved based on predicates.

- Example:

```
SELECT * FROM Employees  
WHERE City = 'x'  
AND Occupation = 'y';
```

Employees Table

ID	Name	Age	City	Occupation
1	Alice	30	Gotham	Data Analyst
2	Bob	22	Metropolis	Grave Digger
3	Charlie	51	Gotham	Product Manager
...	...	...	...	...

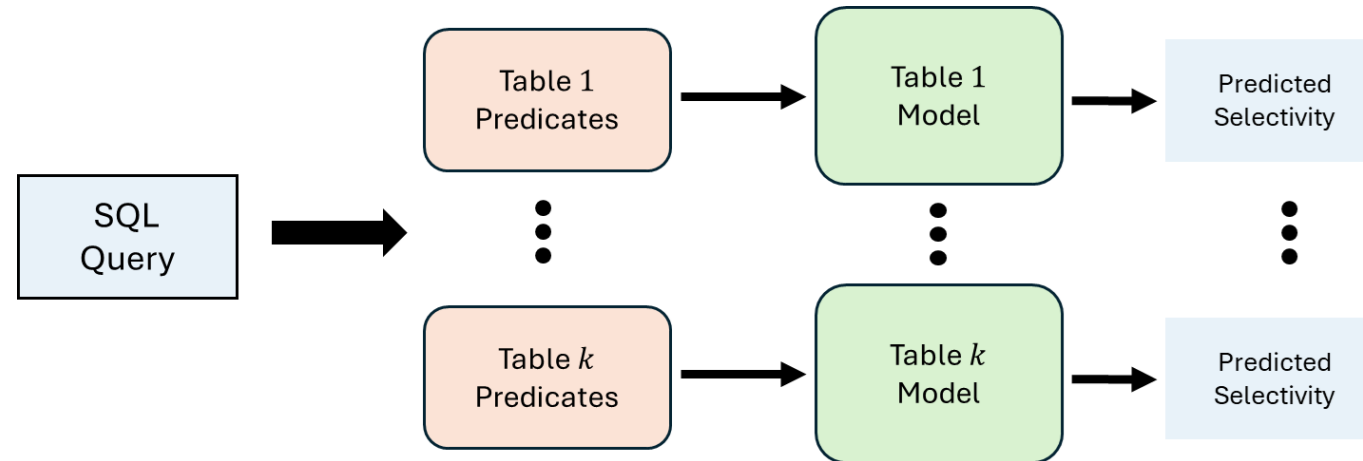
$$\text{selectivity} = P(\text{City} = x, \text{Occupation} = y) \approx P(\text{City} = x) P(\text{Occupation} = y)$$

- **Attribute independence assumption** → often **violated due to correlations**
- **Single attribute distributions**, e.g.  $P(\text{City})$ , from **database catalog statistics** → often **inaccurate/outdated or unavailable**



# Towards Better Selectivity Estimation I: Learned Selectivity Prediction

- **Learning** to perform **selectivity prediction** for **equality/range predicates**.
- A **separate model** for each table.



This approach is powerful because:

- **online** learning → more **accurate/up-to-date statistics**
- can handle **joint selectivities over attributes** → learn **correlations**, no independence assumption
- simple prediction task, simple features → **lightweight** model can be used

# Towards Better Selectivity Estimation II: Learned CDFs

- **Learning** the **cumulative distribution function** (CDF) of table attributes via **monotonic regression**.

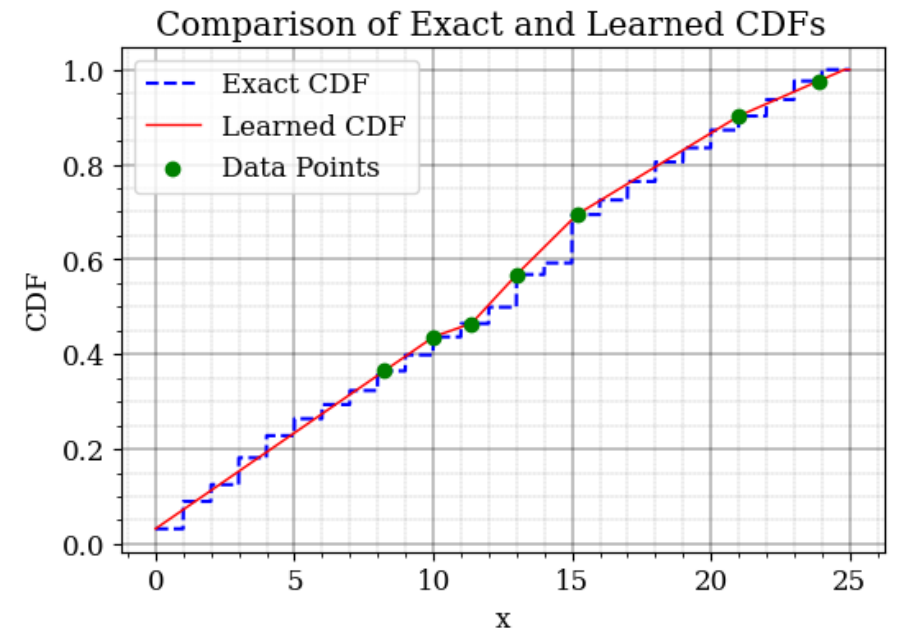
$$CDF_A(x) = P(A \leq x)$$

- **Selectivity can be derived using the CDF**, e.g. for this range predicate  $\{a_1 \leq A \leq a_2\}$ :

$$\text{selectivity} = P(a_1 \leq A \leq a_2) = CDF_A(a_2) - CDF_A(a_1)$$

This approach is powerful because:

- **online** learning
- **high sample efficiency**
- handle **joint selectivities** via **multidimensional monotonic regression**
- **very lightweight**



# Preliminary Results - Benchmarking WFIT vs. MAB

- Experimental environment:
  - Linux Ubuntu OS, 20GB RAM, Intel Core-i7 16 cores, 1TB SSD
  - PostgreSQL + HypoPG extension
- Datasets
  - SSB, TPC-H, TPC-H skew (all SF 10)
- Workload Types - Static, Dynamic Shifting, Dynamic Random
- Query execution with cold cache (OS and PG buffer caches empty at start of each query execution)

## Experiment:

- SSB Dataset
- Static Workload - 10 templates
- 8 rounds

	<b>WFIT + what-if</b>	<b>WFIT + ext.</b>	<b>MAB</b>	<b>No Index</b>
$C_{tot}(W)$	583 s	<b>507 s</b>	1021 s	2258 s
$C_{tr}(W)$	<b>140 s</b>	152 s	455 s	0 s
$C_{exe}(W)$	392 s	<b>355 s</b>	566 s	2258 s

# Conclusions & Future Directions

- WFIT algorithm holds up well against current SOTA MAB!
- External cost model significantly faster (500x) than what-if interface.
- Learned selectivity estimation may be a promising approach.
- However, still have room for improvement:
  - Extend access path enumeration to incorporate multiple index scans for a single table
  - Extend to multi-attribute joint selectivity prediction and multidimensional CDF
  - Scope for speeding up WFIT, certain aspects are “embarrassingly parallel” → multi-process parallelization
- On experimental side:
  - Need to benchmark on larger real-world datasets with highly skewed data, e.g. imdb
  - Use more diverse, complex and larger workloads
  - Design experiments that test ability of learned models to handle joint attributes with high correlation